

Class Challenge: Image Classification of COVID-19 X-rays

Task 1 [Total points: 30]

Setup

- This assignment involves the following packages: 'matplotlib', 'numpy', and 'sklearn'.
- If you are using conda, use the following commands to install the above packages:

```
conda install matplotlib
conda install numpy
conda install -c anaconda scikit-learn
```

- If you are using pip, use the following commands to install the above packages:

```
pip install matplotlib
pip install numpy
pip install sklearn
```

Data

Please download the data using the following link: [COVID-19 \(https://drive.google.com/file/d/1Y88tgqpQ1Pjko_7rntcPowOJs_QNOrJ-/view\)](https://drive.google.com/file/d/1Y88tgqpQ1Pjko_7rntcPowOJs_QNOrJ-/view).

- After downloading 'Covid_Data_GradientCrescent.zip', unzip the file and you should see the following data structure:

```
|--all
|-----train
|-----test
|--two
|-----train
|-----test
```

- Put the 'all' folder, the 'two' folder and this python notebook in the **same directory** so that the following code can correctly locate the data.

[20 points] Binary Classification: COVID-19 vs. Normal

In [12]:

```
import os

import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing.image import ImageDataGenerator

os.environ['OMP_NUM_THREADS'] = '1'
os.environ['CUDA_VISIBLE_DEVICES'] = '-1'
tf.__version__
```

Out[12]:

'2.3.1'

Load Image Data

In [13]:

```
DATA_LIST = os.listdir('two/train')
DATASET_PATH = 'two/train'
TEST_DIR = 'two/test'
IMAGE_SIZE = (224, 224)
NUM_CLASSES = len(DATA_LIST)
BATCH_SIZE = 10 # try reducing batch size or freeze more layers if your GPU runs out of memory
NUM_EPOCHS = 40
LEARNING_RATE = 0.0005 # start off with high rate first 0.001 and experiment with reducing it gradually
```

Generate Training and Validation Batches

In [14]:

```
train_datagen = ImageDataGenerator(rescale=1./255,rotation_range=50,featurewise_center = True,
                                   featurewise_std_normalization = True,width_shift_range=0.2,
                                   height_shift_range=0.2,shear_range=0.25,zoom_range=0.1,
                                   zca_whitening = True,channel_shift_range = 20,
                                   horizontal_flip = True,vertical_flip = True,
                                   validation_split = 0.2,fill_mode='constant')

train_batches = train_datagen.flow_from_directory(DATASET_PATH,target_size=IMAGE_SIZE,
                                                  shuffle=True,batch_size=BATCH_SIZE,
                                                  subset = "training",seed=42,
                                                  class_mode="binary")

valid_batches = train_datagen.flow_from_directory(DATASET_PATH,target_size=IMAGE_SIZE,
                                                  shuffle=True,batch_size=BATCH_SIZE,
                                                  subset = "validation",seed=42,
                                                  class_mode="binary")
```

Found 104 images belonging to 2 classes.
Found 26 images belonging to 2 classes.

[10 points] Build Model

Hint: Starting from a pre-trained model typically helps performance on a new task, e.g. starting with weights obtained by training on ImageNet.

In [15]:

```
vgg16 = tf.keras.applications.VGG16(weights='imagenet', include_top=False, input_shape = (224, 224, 3))
vgg19 = tf.keras.applications.VGG19(include_top=False, weights='imagenet', input_shape=(224, 224, 3))
model = tf.keras.models.Sequential([
    vgg16,
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(256, activation='relu', name='dense_feature'),
    tf.keras.layers.Dropout(0.25),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
vgg16.trainable=False
model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 7, 7, 512)	14714688
flatten_3 (Flatten)	(None, 25088)	0
dense_feature (Dense)	(None, 256)	6422784
dropout_1 (Dropout)	(None, 256)	0
dense_3 (Dense)	(None, 1)	257
Total params: 21,137,729		
Trainable params: 6,423,041		
Non-trainable params: 14,714,688		

[5 points] Train Model

In [16]:

```
#FIT MODEL
print(len(train_batches))
print(len(valid_batches))

STEP_SIZE_TRAIN=train_batches.n//train_batches.batch_size
STEP_SIZE_VALID=valid_batches.n//valid_batches.batch_size

opt = tf.keras.optimizers.Adam(learning_rate=LEARNING_RATE)
model.compile(optimizer=opt,loss=tf.keras.losses.BinaryCrossentropy(),metrics=['accuracy'])
history = model.fit(train_batches, epochs=NUM_EPOCHS, validation_data=valid_batches,batch_size=5
, steps_per_epoch=STEP_SIZE_TRAIN, validation_steps = STEP_SIZE_VALID)
```

11

3

Epoch 1/40

10/10 [=====] - 13s 1s/step - loss: 1.6453 - accuracy: 0.5426 - val_loss: 0.2575 - val_accuracy: 0.9000

Epoch 2/40

10/10 [=====] - 12s 1s/step - loss: 0.7682 - accuracy: 0.7021 - val_loss: 0.3129 - val_accuracy: 0.8000

Epoch 3/40

10/10 [=====] - 12s 1s/step - loss: 0.4550 - accuracy: 0.8298 - val_loss: 0.1807 - val_accuracy: 0.9500

Epoch 4/40

10/10 [=====] - 13s 1s/step - loss: 0.3345 - accuracy: 0.8600 - val_loss: 0.2312 - val_accuracy: 0.9000

Epoch 5/40

10/10 [=====] - 13s 1s/step - loss: 0.2664 - accuracy: 0.8900 - val_loss: 0.2187 - val_accuracy: 0.8500

Epoch 6/40

10/10 [=====] - 13s 1s/step - loss: 0.2515 - accuracy: 0.8936 - val_loss: 0.1000 - val_accuracy: 1.0000

Epoch 7/40

10/10 [=====] - 13s 1s/step - loss: 0.3009 - accuracy: 0.9043 - val_loss: 0.1053 - val_accuracy: 1.0000

Epoch 8/40

10/10 [=====] - 12s 1s/step - loss: 0.3481 - accuracy: 0.8936 - val_loss: 0.0760 - val_accuracy: 1.0000

Epoch 9/40

10/10 [=====] - 13s 1s/step - loss: 0.2886 - accuracy: 0.8723 - val_loss: 0.0694 - val_accuracy: 1.0000

Epoch 10/40

10/10 [=====] - 13s 1s/step - loss: 0.1577 - accuracy: 0.9255 - val_loss: 0.0628 - val_accuracy: 1.0000

Epoch 11/40

10/10 [=====] - 14s 1s/step - loss: 0.1261 - accuracy: 0.9362 - val_loss: 0.0998 - val_accuracy: 0.9500

Epoch 12/40

10/10 [=====] - 13s 1s/step - loss: 0.1626 - accuracy: 0.9468 - val_loss: 0.0385 - val_accuracy: 1.0000

Epoch 13/40

10/10 [=====] - 13s 1s/step - loss: 0.1440 - accuracy: 0.9362 - val_loss: 0.0176 - val_accuracy: 1.0000

Epoch 14/40

10/10 [=====] - 12s 1s/step - loss: 0.1196 - accuracy: 0.9681 - val_loss: 0.0519 - val_accuracy: 1.0000

Epoch 15/40

10/10 [=====] - 12s 1s/step - loss: 0.1579 - accuracy: 0.9255 - val_loss: 0.0768 - val_accuracy: 0.9500

Epoch 16/40

10/10 [=====] - 13s 1s/step - loss: 0.1992 - accuracy: 0.9200 - val_loss: 0.0468 - val_accuracy: 1.0000

Epoch 17/40

10/10 [=====] - 15s 1s/step - loss: 0.1541 - accuracy: 0.9500 - val_loss: 0.0874 - val_accuracy: 1.0000

Epoch 18/40

10/10 [=====] - 14s 1s/step - loss: 0.1677 - accuracy: 0.9255 - val_loss: 0.0506 - val_accuracy: 1.0000

Epoch 19/40

10/10 [=====] - 13s 1s/step - loss: 0.1508 - accuracy: 0.9362 - val_loss: 0.0362 - val_accuracy: 1.0000

Epoch 20/40

10/10 [=====] - 13s 1s/step - loss: 0.1135 - accuracy: 0.

```
9574 - val_loss: 0.0205 - val_accuracy: 1.0000
Epoch 21/40
10/10 [=====] - 13s 1s/step - loss: 0.1064 - accuracy: 0.
9681 - val_loss: 0.3148 - val_accuracy: 0.8500
Epoch 22/40
10/10 [=====] - 15s 2s/step - loss: 0.1326 - accuracy: 0.
9500 - val_loss: 0.0344 - val_accuracy: 1.0000
Epoch 23/40
10/10 [=====] - 17s 2s/step - loss: 0.0905 - accuracy: 0.
9681 - val_loss: 0.0667 - val_accuracy: 0.9500
Epoch 24/40
10/10 [=====] - 13s 1s/step - loss: 0.1231 - accuracy: 0.
9574 - val_loss: 0.0382 - val_accuracy: 1.0000
Epoch 25/40
10/10 [=====] - 13s 1s/step - loss: 0.0957 - accuracy: 0.
9681 - val_loss: 0.2405 - val_accuracy: 0.9500
Epoch 26/40
10/10 [=====] - 13s 1s/step - loss: 0.1683 - accuracy: 0.
9362 - val_loss: 0.0198 - val_accuracy: 1.0000
Epoch 27/40
10/10 [=====] - 12s 1s/step - loss: 0.2280 - accuracy: 0.
8936 - val_loss: 0.1361 - val_accuracy: 0.9500
Epoch 28/40
10/10 [=====] - 13s 1s/step - loss: 0.0878 - accuracy: 0.
9681 - val_loss: 0.0562 - val_accuracy: 1.0000
Epoch 29/40
10/10 [=====] - 13s 1s/step - loss: 0.2063 - accuracy: 0.
9043 - val_loss: 0.0728 - val_accuracy: 1.0000
Epoch 30/40
10/10 [=====] - 12s 1s/step - loss: 0.2714 - accuracy: 0.
8936 - val_loss: 0.4279 - val_accuracy: 0.8000
Epoch 31/40
10/10 [=====] - 14s 1s/step - loss: 0.1848 - accuracy: 0.
9362 - val_loss: 0.1377 - val_accuracy: 0.9000
Epoch 32/40
10/10 [=====] - 15s 1s/step - loss: 0.1480 - accuracy: 0.
9400 - val_loss: 0.1649 - val_accuracy: 0.9500
Epoch 33/40
10/10 [=====] - 13s 1s/step - loss: 0.0619 - accuracy: 0.
9787 - val_loss: 0.2234 - val_accuracy: 0.9000
Epoch 34/40
10/10 [=====] - 12s 1s/step - loss: 0.1411 - accuracy: 0.
9255 - val_loss: 0.0192 - val_accuracy: 1.0000
Epoch 35/40
10/10 [=====] - 12s 1s/step - loss: 0.1318 - accuracy: 0.
9468 - val_loss: 0.2069 - val_accuracy: 0.9500
Epoch 36/40
10/10 [=====] - 12s 1s/step - loss: 0.1266 - accuracy: 0.
9681 - val_loss: 0.0521 - val_accuracy: 1.0000
Epoch 37/40
10/10 [=====] - 13s 1s/step - loss: 0.0601 - accuracy: 0.
9787 - val_loss: 0.2902 - val_accuracy: 0.9000
Epoch 38/40
10/10 [=====] - 12s 1s/step - loss: 0.0619 - accuracy: 0.
9894 - val_loss: 0.1917 - val_accuracy: 0.9000
Epoch 39/40
10/10 [=====] - 13s 1s/step - loss: 0.1160 - accuracy: 0.
9574 - val_loss: 0.0924 - val_accuracy: 0.9500
Epoch 40/40
10/10 [=====] - 13s 1s/step - loss: 0.0860 - accuracy: 0.
9500 - val_loss: 0.1260 - val_accuracy: 0.9500
```



[5 points] Plot Accuracy and Loss During Training

In [10]:

```
import matplotlib.pyplot as plt

plt.title('Accuracy over 40 Epochs')
plt.plot(history.history['accuracy'], label='Train_acc')
plt.plot(history.history['val_accuracy'], label='Test_acc')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

plt.title('Loss over 40 Epochs')
plt.plot(history.history['loss'], label='Train_loss')
plt.plot(history.history['val_loss'], label='Test_loss')
plt.legend(loc='upper right')
plt.grid(True)
plt.show()
```

**Plot Test Results**

In [11]:

```
import matplotlib.image as mpimg

test_datagen = ImageDataGenerator(rescale=1. / 255)
eval_generator = test_datagen.flow_from_directory(TEST_DIR, target_size=IMAGE_SIZE,
                                                  batch_size=1, shuffle=True, seed=42, class_mode=
"binary")
eval_generator.reset()
pred = model.predict_generator(eval_generator, 18, verbose=1)
for index, probability in enumerate(pred):
    image_path = TEST_DIR + "/" + eval_generator.filenames[index]
    image = mpimg.imread(image_path)
    if image.ndim < 3:
        image = np.reshape(image, (image.shape[0], image.shape[1], 1))
        image = np.concatenate([image, image, image], 2)
    #     print(image.shape)

    pixels = np.array(image)
    plt.imshow(pixels)

    print(eval_generator.filenames[index])
    if probability > 0.5:
        plt.title("%.2f" % (probability[0]*100) + "% Normal")
    else:
        plt.title("%.2f" % ((1-probability[0])*100) + "% COVID19 Pneumonia")
    plt.show()
```


Found 18 images belonging to 2 classes.


WARNING:tensorflow:From <ipython-input-11-543347a5fba8>:7: Model.predict_generator (from tensorflow.python.keras.engine.training) is deprecated and will be removed in a future version.

Instructions for updating:


Please use Model.predict, which supports generators.

18/18 [=====] - 2s 111ms/step


covidWnejmoa2001191_f3-PA.jpeg




covidWnejmoa2001191_f4.jpeg




covidWnejmoa2001191_f5-PA.jpeg




covidWradiol.2020200490.fig3.jpeg




covidWryct.2020200028.fig1a.jpeg




covidWryct.2020200034.fig2.jpeg




covidWryct.2020200034.fig5-day0.jpeg




covidWryct.2020200034.fig5-day4.jpeg




covidWryct.2020200034.fig5-day7.jpeg




normal\WNORMAL2-IM-1385-0001.jpeg




normal\WNORMAL2-IM-1396-0001.jpeg




normal\WNORMAL2-IM-1400-0001.jpeg




normal\WNORMAL2-IM-1401-0001.jpeg




normal\WNORMAL2-IM-1406-0001.jpeg




normal\WNORMAL2-IM-1412-0001.jpeg




normal\WNORMAL2-IM-1419-0001.jpeg



normal\WNORMAL2-IM-1422-0001.jpeg



normal\WNORMAL2-IM-1423-0001.jpeg



[10 points] TSNE Plot

t-Distributed Stochastic Neighbor Embedding (t-SNE) is a widely used technique for dimensionality reduction that is particularly well suited for the visualization of high-dimensional datasets. After training is complete, extract features from a specific deep layer of your choice, use t-SNE to reduce the dimensionality of your extracted features to 2 dimensions and plot the resulting 2D features.

In [13]:

```
from sklearn.manifold import TSNE

intermediate_layer_model = models.Model(inputs=model.input,
                                         outputs=model.get_layer('dense_feature').output)
tsne_data_generator = test_datagen.flow_from_directory(DATASET_PATH, target_size=IMAGE_SIZE,
                                                       batch_size=1, shuffle=True, seed=42, class_mode=
"binary")

labels = []
num = len(tsne_data_generator)
for i in range(num):
    labels.extend(np.array(tsne_data_generator[i][1]))

#feature extraction
tsne_data_generator.reset()
features = intermediate_layer_model.predict_generator(tsne_data_generator)

#compress the dimensionality
tsne = TSNE(n_components=2)
tsne_result = tsne.fit_transform(features)
x=[i[0] for i in tsne_result]
y=[i[1] for i in tsne_result]

#plotting values
covid_x = []
covid_y = []
non_covid_x = []
non_covid_y = []
for i in range(len(labels)):
    if labels[i] == 0.0:
        non_covid_x.append(x[i])
        non_covid_y.append(y[i])
    else:
        covid_x.append(x[i])
        covid_y.append(y[i])

#plotting
plt.scatter(non_covid_x, non_covid_y, label='Normal')
plt.scatter(covid_x, covid_y, label='COVID-19')
plt.legend(loc='upper left')
plt.show()
```


Found 130 images belonging to 2 classes.

130

