

# Class Challenge: Image Classification of COVID-19 X-rays

## Task 1 [Total points: 30]

### Setup

- This assignment involves the following packages: 'matplotlib', 'numpy', and 'sklearn'.
- If you are using conda, use the following commands to install the above packages:

```
conda install matplotlib
conda install numpy
conda install -c anaconda scikit-learn
```

- If you are using pip, use the following commands to install the above packages:

```
pip install matplotlib
pip install numpy
pip install sklearn
```

### Data

Please download the data using the following link: [COVID-19 \(https://drive.google.com/file/d/1Y88tgqpQ1Pjko\\_7rntcPowOJs\\_QNOrJ-/view\)](https://drive.google.com/file/d/1Y88tgqpQ1Pjko_7rntcPowOJs_QNOrJ-/view).

- After downloading 'Covid\_Data\_GradientCrescent.zip', unzip the file and you should see the following data structure:

```
|--all
|-----train
|-----test
|--two
|-----train
|-----test
```

- Put the 'all' folder, the 'two' folder and this python notebook in the **same directory** so that the following code can correctly locate the data.

## [20 points] Binary Classification: COVID-19 vs. Normal

In [1]:

```
import os

import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing.image import ImageDataGenerator

os.environ['OMP_NUM_THREADS'] = '1'
os.environ['CUDA_VISIBLE_DEVICES'] = '-1'
tf.__version__
```

Out[1]:

'2.3.1'

## Load Image Data

In [2]:

```
DATA_LIST = os.listdir('two/train')
DATASET_PATH = 'two/train'
TEST_DIR = 'two/test'
IMAGE_SIZE = (224, 224)
NUM_CLASSES = len(DATA_LIST)
BATCH_SIZE = 10 # try reducing batch size or freeze more layers if your GPU runs out of memory
NUM_EPOCHS = 40
LEARNING_RATE = 0.0005 # start off with high rate first 0.001 and experiment with reducing it gradually
```

## Generate Training and Validation Batches

In [3]:

```

train_datagen = ImageDataGenerator(rescale=1./255,rotation_range=50,featurewise_center = True,
                                   featurewise_std_normalization = True,width_shift_range=0.2,
                                   height_shift_range=0.2,shear_range=0.25,zoom_range=0.1,
                                   zca_whitening = True,channel_shift_range = 20,
                                   horizontal_flip = True,vertical_flip = True,
                                   validation_split = 0.2,fill_mode='constant')

train_batches = train_datagen.flow_from_directory(DATASET_PATH,target_size=IMAGE_SIZE,
                                                  shuffle=False,batch_size=BATCH_SIZE,
                                                  subset = "training",seed=42,
                                                  class_mode="binary")

valid_batches = train_datagen.flow_from_directory(DATASET_PATH,target_size=IMAGE_SIZE,
                                                  shuffle=False,batch_size=BATCH_SIZE,
                                                  subset = "validation",seed=42,
                                                  class_mode="binary")

```

C:\Users\WSoon\Whan Park\Anaconda3\lib\site-packages\keras\_preprocessing\image\image\_data\_generator.py:342: UserWarning: This ImageDataGenerator specifies `zca\_whitening` which overrides setting of `featurewise\_std\_normalization`.  
 warnings.warn('This ImageDataGenerator specifies '

Found 104 images belonging to 2 classes.

Found 26 images belonging to 2 classes.

### [10 points] Build Model

Hint: Starting from a pre-trained model typically helps performance on a new task, e.g. starting with weights obtained by training on ImageNet.

In [4]:

```

vgg16 = tf.keras.applications.VGG16(weights='imagenet', include_top=False, input_shape = (224, 224, 3))
#vgg19 = tf.keras.applications.VGG19(include_top=False, weights='imagenet', input_shape=(224, 224, 3))
model = tf.keras.models.Sequential([
    vgg16,
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(256, activation='relu', name='dense_feature'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
vgg16.trainable=False
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 7, 7, 512)	14714688
flatten (Flatten)	(None, 25088)	0
dense_feature (Dense)	(None, 256)	6422784
dropout (Dropout)	(None, 256)	0
dense (Dense)	(None, 1)	257
Total params: 21,137,729		
Trainable params: 6,423,041		
Non-trainable params: 14,714,688		

**[5 points] Train Model**

In [5]:

```
#FIT MODEL
print(len(train_batches))
print(len(valid_batches))

STEP_SIZE_TRAIN=train_batches.n//train_batches.batch_size
STEP_SIZE_VALID=valid_batches.n//valid_batches.batch_size

opt = tf.keras.optimizers.Adam(learning_rate=LEARNING_RATE)
model.compile(optimizer=opt,loss=tf.keras.losses.BinaryCrossentropy(),metrics=['accuracy'])
history = model.fit(train_batches, epochs=NUM_EPOCHS, validation_data=valid_batches,batch_size=5
, steps_per_epoch=STEP_SIZE_TRAIN, validation_steps = STEP_SIZE_VALID)
```

11  
3

```
C:\Users\WSo Whan Park\Anaconda3\lib\site-packages\keras_preprocessing\image\image_data_generator.py:720: UserWarning: This ImageDataGenerator specifies `featurewise_center`, but it hasn't been fit on any training data. Fit it first by calling `.fit(numpy_data)`.
```

```
warnings.warn('This ImageDataGenerator specifies '
```

```
C:\Users\WSo Whan Park\Anaconda3\lib\site-packages\keras_preprocessing\image\image_data_generator.py:739: UserWarning: This ImageDataGenerator specifies `zca_whitening`, but it hasn't been fit on any training data. Fit it first by calling `.fit(numpy_data)`.
```

```
warnings.warn('This ImageDataGenerator specifies '
```

Epoch 1/40  
10/10 [=====] - 12s 1s/step - loss: 4.3465 - accuracy: 0.5957 - val\_loss: 1.2977 - val\_accuracy: 0.6000

Epoch 2/40  
10/10 [=====] - 10s 1s/step - loss: 1.1545 - accuracy: 0.4787 - val\_loss: 1.1984 - val\_accuracy: 0.4000

Epoch 3/40  
10/10 [=====] - 11s 1s/step - loss: 0.9066 - accuracy: 0.4468 - val\_loss: 0.5460 - val\_accuracy: 0.9000

Epoch 4/40  
10/10 [=====] - 11s 1s/step - loss: 0.6321 - accuracy: 0.6702 - val\_loss: 0.5080 - val\_accuracy: 0.9000

Epoch 5/40  
10/10 [=====] - 12s 1s/step - loss: 0.6412 - accuracy: 0.6277 - val\_loss: 0.6674 - val\_accuracy: 0.4500

Epoch 6/40  
10/10 [=====] - 12s 1s/step - loss: 0.6128 - accuracy: 0.6277 - val\_loss: 0.4855 - val\_accuracy: 0.9000

Epoch 7/40  
10/10 [=====] - 12s 1s/step - loss: 0.6052 - accuracy: 0.6596 - val\_loss: 0.5414 - val\_accuracy: 0.6500

Epoch 8/40  
10/10 [=====] - 12s 1s/step - loss: 0.5548 - accuracy: 0.7021 - val\_loss: 0.4583 - val\_accuracy: 0.9000

Epoch 9/40  
10/10 [=====] - 12s 1s/step - loss: 0.5943 - accuracy: 0.7447 - val\_loss: 0.4474 - val\_accuracy: 1.0000

Epoch 10/40  
10/10 [=====] - 12s 1s/step - loss: 0.5415 - accuracy: 0.7900 - val\_loss: 0.4614 - val\_accuracy: 0.7500

Epoch 11/40  
10/10 [=====] - 12s 1s/step - loss: 0.4608 - accuracy: 0.7979 - val\_loss: 0.4240 - val\_accuracy: 0.8500

Epoch 12/40  
10/10 [=====] - 12s 1s/step - loss: 0.4790 - accuracy: 0.7766 - val\_loss: 0.3783 - val\_accuracy: 1.0000

Epoch 13/40  
10/10 [=====] - 13s 1s/step - loss: 0.4838 - accuracy: 0.8298 - val\_loss: 0.3577 - val\_accuracy: 0.9500

Epoch 14/40  
10/10 [=====] - 14s 1s/step - loss: 0.4557 - accuracy: 0.7660 - val\_loss: 0.3831 - val\_accuracy: 0.9000

Epoch 15/40  
10/10 [=====] - 14s 1s/step - loss: 0.4448 - accuracy: 0.8700 - val\_loss: 0.3329 - val\_accuracy: 1.0000

Epoch 16/40  
10/10 [=====] - 13s 1s/step - loss: 0.4430 - accuracy: 0.8511 - val\_loss: 0.3572 - val\_accuracy: 0.9000

Epoch 17/40  
10/10 [=====] - 14s 1s/step - loss: 0.4411 - accuracy: 0.8191 - val\_loss: 0.3630 - val\_accuracy: 0.8500

Epoch 18/40  
10/10 [=====] - 13s 1s/step - loss: 0.3972 - accuracy: 0.7900 - val\_loss: 0.3934 - val\_accuracy: 0.8500

Epoch 19/40  
10/10 [=====] - 13s 1s/step - loss: 0.3557 - accuracy: 0.9100 - val\_loss: 0.3131 - val\_accuracy: 0.9000

Epoch 20/40  
10/10 [=====] - 13s 1s/step - loss: 0.4550 - accuracy: 0.7766 - val\_loss: 0.3457 - val\_accuracy: 0.9000

Epoch 21/40

```
10/10 [=====] - 13s 1s/step - loss: 0.3790 - accuracy: 0.8617 - val_loss: 0.2498 - val_accuracy: 1.0000
Epoch 22/40
10/10 [=====] - 12s 1s/step - loss: 0.3942 - accuracy: 0.8298 - val_loss: 0.2830 - val_accuracy: 0.9500
Epoch 23/40
10/10 [=====] - 13s 1s/step - loss: 0.4161 - accuracy: 0.8298 - val_loss: 0.3130 - val_accuracy: 0.9500
Epoch 24/40
10/10 [=====] - 13s 1s/step - loss: 0.3679 - accuracy: 0.8617 - val_loss: 0.2419 - val_accuracy: 0.9500
Epoch 25/40
10/10 [=====] - 13s 1s/step - loss: 0.3459 - accuracy: 0.8298 - val_loss: 0.3185 - val_accuracy: 0.9000
Epoch 26/40
10/10 [=====] - 12s 1s/step - loss: 0.4319 - accuracy: 0.8085 - val_loss: 0.2662 - val_accuracy: 0.9000
Epoch 27/40
10/10 [=====] - 14s 1s/step - loss: 0.4049 - accuracy: 0.8191 - val_loss: 0.3085 - val_accuracy: 0.9000
Epoch 28/40
10/10 [=====] - 13s 1s/step - loss: 0.2970 - accuracy: 0.9255 - val_loss: 0.2134 - val_accuracy: 0.9500
Epoch 29/40
10/10 [=====] - 13s 1s/step - loss: 0.3631 - accuracy: 0.8511 - val_loss: 0.2216 - val_accuracy: 1.0000
Epoch 30/40
10/10 [=====] - 12s 1s/step - loss: 0.3764 - accuracy: 0.8511 - val_loss: 0.2362 - val_accuracy: 1.0000
Epoch 31/40
10/10 [=====] - 13s 1s/step - loss: 0.2534 - accuracy: 0.9043 - val_loss: 0.2574 - val_accuracy: 0.9000
Epoch 32/40
10/10 [=====] - 13s 1s/step - loss: 0.2834 - accuracy: 0.9362 - val_loss: 0.1635 - val_accuracy: 0.9500
Epoch 33/40
10/10 [=====] - 13s 1s/step - loss: 0.2673 - accuracy: 0.8936 - val_loss: 0.3480 - val_accuracy: 0.9500
Epoch 34/40
10/10 [=====] - 12s 1s/step - loss: 0.3355 - accuracy: 0.8191 - val_loss: 0.2498 - val_accuracy: 0.8500
Epoch 35/40
10/10 [=====] - 12s 1s/step - loss: 0.2896 - accuracy: 0.8617 - val_loss: 0.1762 - val_accuracy: 1.0000
Epoch 36/40
10/10 [=====] - 13s 1s/step - loss: 0.2355 - accuracy: 0.9500 - val_loss: 0.1988 - val_accuracy: 1.0000
Epoch 37/40
10/10 [=====] - 12s 1s/step - loss: 0.3001 - accuracy: 0.8936 - val_loss: 0.2350 - val_accuracy: 0.9500
Epoch 38/40
10/10 [=====] - 13s 1s/step - loss: 0.3193 - accuracy: 0.8500 - val_loss: 0.1888 - val_accuracy: 0.9500
Epoch 39/40
10/10 [=====] - 12s 1s/step - loss: 0.2951 - accuracy: 0.8830 - val_loss: 0.1466 - val_accuracy: 1.0000
Epoch 40/40
10/10 [=====] - 12s 1s/step - loss: 0.3108 - accuracy: 0.8511 - val_loss: 0.2260 - val_accuracy: 0.9500
```





**[5 points] Plot Accuracy and Loss During Training**

In [6]:

```
import matplotlib.pyplot as plt

plt.title('Accuracy over 40 Epochs')
plt.plot(history.history['accuracy'], label='Train_acc')
plt.plot(history.history['val_accuracy'], label='Test_acc')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

plt.title('Loss over 40 Epochs')
plt.plot(history.history['loss'], label='Train_loss')
plt.plot(history.history['val_loss'], label='Test_loss')
plt.legend(loc='upper right')
plt.grid(True)
plt.show()
```

**Plot Test Results**

In [7]:

```
import matplotlib.image as mpimg

test_datagen = ImageDataGenerator(rescale=1. / 255)
eval_generator = test_datagen.flow_from_directory(TEST_DIR, target_size=IMAGE_SIZE,
                                                  batch_size=1, shuffle=False, seed=42, class_mode
                                                  ="binary")
eval_generator.reset()
pred = model.predict_generator(eval_generator, 18, verbose=1)
for index, probability in enumerate(pred):
    image_path = TEST_DIR + "/" + eval_generator.filesnames[index]
    image = mpimg.imread(image_path)
    if image.ndim < 3:
        image = np.reshape(image, (image.shape[0], image.shape[1], 1))
        image = np.concatenate([image, image, image], 2)
    #     print(image.shape)

    pixels = np.array(image)
    plt.imshow(pixels)

    print(eval_generator.filesnames[index])
    if probability > 0.5:
        plt.title("%.2f" % (probability[0]*100) + "% Normal")
    else:
        plt.title("%.2f" % ((1-probability[0])*100) + "% COVID19 Pneumonia")
    plt.show()
```

Found 18 images belonging to 2 classes.


WARNING:tensorflow:From <ipython-input-7-aa0cc9a8f179>:7: Model.predict\_generator (from tensorflow.python.keras.engine.training) is deprecated and will be removed in a future version.

Instructions for updating:


Please use Model.predict, which supports generators.

18/18 [=====] - 2s 123ms/step


covidWnejmoa2001191\_f3-PA.jpeg




covidWnejmoa2001191\_f4.jpeg




covidWnejmoa2001191\_f5-PA.jpeg




covidWradiol.2020200490.fig3.jpeg




covidWryct.2020200028.fig1a.jpeg




covidWryct.2020200034.fig2.jpeg




covidWryct.2020200034.fig5-day0.jpeg




covidWryct.2020200034.fig5-day4.jpeg




covidWryct.2020200034.fig5-day7.jpeg




normalWNORMAL2-IM-1385-0001.jpeg



normalWNORMAL2-IM-1396-0001.jpeg




normal\WNORMAL2-IM-1400-0001.jpeg




normal\WNORMAL2-IM-1401-0001.jpeg







normal\WNORMAL2-IM-1406-0001.jpeg




normal\WNORMAL2-IM-1412-0001.jpeg




normal\WNORMAL2-IM-1419-0001.jpeg



normal\WNORMAL2-IM-1422-0001.jpeg



normal\WNORMAL2-IM-1423-0001.jpeg



In [12]:

```
print('Test accuracy:', 17/18)
```

Test accuracy: 0.9444444444444444

## [10 points] TSNE Plot

t-Distributed Stochastic Neighbor Embedding (t-SNE) is a widely used technique for dimensionality reduction that is particularly well suited for the visualization of high-dimensional datasets. After training is complete, extract features from a specific deep layer of your choice, use t-SNE to reduce the dimensionality of your extracted features to 2 dimensions and plot the resulting 2D features.

In [8]:

```
from sklearn.manifold import TSNE

intermediate_layer_model = tf.keras.models.Model(inputs=model.input,
                                                  outputs=model.get_layer('dense_feature').output)
tsne_data_generator = test_datagen.flow_from_directory(DATASET_PATH, target_size=IMAGE_SIZE,
                                                       batch_size=1, shuffle=False, seed=42, class_mode
                                                       ="binary")

labels = []
num = len(tsne_data_generator)
for i in range(num):
    labels.extend(np.array(tsne_data_generator[i][1]))

#feature extraction
tsne_data_generator.reset()
features = intermediate_layer_model.predict_generator(tsne_data_generator)

#compress the dimensionality
tsne = TSNE(n_components=2)
tsne_result = tsne.fit_transform(features)
x=[i[0] for i in tsne_result]
y=[i[1] for i in tsne_result]

#plotting values
covid_x = []
covid_y = []
non_covid_x = []
non_covid_y = []
for i in range(len(labels)):
    if labels[i] == 0.0:
        non_covid_x.append(x[i])
        non_covid_y.append(y[i])
    else:
        covid_x.append(x[i])
        covid_y.append(y[i])

#plotting
plt.scatter(non_covid_x, non_covid_y, label='Normal')
plt.scatter(covid_x, covid_y, label='COVID-19')
plt.legend(loc='upper left')
plt.show()
```

Found 130 images belonging to 2 classes.

